

## Lesson 02 Demo 03

### Launching a Pod and Establishing an Associated Service

**Objective:** To integrate Kubernetes deployments with services to achieve scalable and accessible pod configurations

**Tools required:** kubeadm, kubectl, kubelet, and containerd

**Prerequisites:** A Kubernetes cluster (refer to Demo 01 from Lesson 01 for setting up a cluster)

Steps to be followed:

1. Create a deployment object
2. Create a service with a label selector for deployment

#### Step 1: Create a deployment object

- 1.1 Create a YAML file for the deployment using the following command:  
**vi mydeployment.yaml**

```
labsuser@master:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
master.example.com                  Ready    control-plane   10d   v1.30.5
worker-node-1.example.com           Ready    <none>         10d   v1.30.4
worker-node-2.example.com           Ready    <none>         10d   v1.30.4
labsuser@master:~$ vi mydeployment.yaml
```

1.2 Add the following code to **mydeployment.yaml**:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: httpd
  replicas: 2
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
      - name: httpd
        image: httpd:latest
        ports:
        - containerPort: 80
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: httpd
  replicas: 2
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
      - name: httpd
        image: httpd:latest
        ports:
        - containerPort: 80
```

1.3 Apply the deployment object using the following command:

**kubectl apply -f mydeployment.yaml**

```
labsuser@master:~$ vi mydeployment.yaml
labsuser@master:~$ kubectl apply -f mydeployment.yaml
deployment.apps/nginx-deployment created
labsuser@master:~$
```

1.4 Verify the deployment and its pods using the following commands:

**kubectl get deployment**

**kubectl get pods**

```
labsuser@master:~$ kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
admin         0/1     1            0           5d22h
nginx         1/1     1            1           4d4h
nginx-deployment 2/2     2            2           23s
labsuser@master:~$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
admin-56d684dff9-zjfhc             0/1     ImagePullBackOff    0          5d22h
counter                             1/1     Running             4 (27m ago) 5d2h
nginx-7854ff8877-mvrtr             1/1     Running             1 (4d ago)  4d4h
nginx-deployment-6d6b866d8f-bw8xr  1/1     Running             0          38s
nginx-deployment-6d6b866d8f-r7pnj  1/1     Running             0          38s
pod-demo                           1/1     Running             8 (27m ago) 9d
labsuser@master:~$
```

## Step 2: Create a service with a label selector for deployment

2.1 Create a new YAML file for the service using the command below:

**vi myservice.yaml**

```
labsuser@master:~$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
admin-56d684dff9-zjfhc             0/1     ImagePullBackOff    0          5d22h
counter                             1/1     Running             4 (27m ago) 5d2h
nginx-7854ff8877-mvrtr             1/1     Running             1 (4d ago)  4d4h
nginx-deployment-6d6b866d8f-bw8xr  1/1     Running             0          38s
nginx-deployment-6d6b866d8f-r7pnj  1/1     Running             0          38s
pod-demo                           1/1     Running             8 (27m ago) 9d
labsuser@master:~$ vi myservice.yaml
```

2.2 Add the following code to **myservice.yaml**:

```
apiVersion: v1
kind: Service
metadata:
  name: myservice
spec:
  selector:
    app: httpd
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: myservice
spec:
  selector:
    app: httpd
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

2.3 Apply the service object using the following command:

```
kubectl apply -f myservice.yaml
```

```
labsuser@master:~$ vi myservice.yaml
labsuser@master:~$ kubectl apply -f myservice.yaml
service/myservice created
labsuser@master:~$
```

- 2.4 Describe the service to verify its connection to the pods using the command below:  
**kubectl describe svc myservice**

```
labsuser@master:~$ kubectl apply -f myservice.yaml
service/myservice created
labsuser@master:~$ kubectl describe svc myservice
Name:                 myservice
Namespace:            default
Labels:               <none>
Annotations:          <none>
Selector:             app=httpd
Type:                 ClusterIP
IP Family Policy:     SingleStack
IP Families:          IPv4
IP:                   10.99.56.93
IPs:                  10.99.56.93
Port:                 <unset> 8080/TCP
TargetPort:           80/TCP
Endpoints:            192.168.232.208:80,192.168.47.146:80
Session Affinity:     None
Events:              <none>
labsuser@master:~$
```

- 2.5 Check the targeted pods by listing them using the service's selector using the command below:  
**kubectl get pods -l app=httpd**

```
labsuser@master:~$ kubectl get pods -l app=httpd
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-6d6b866d8f-bw8xr   1/1     Running   0           4m55s
nginx-deployment-6d6b866d8f-r7pnj   1/1     Running   0           4m55s
labsuser@master:~$
```

2.6 List the service's endpoints to view the IP addresses of the pods it targets:

**kubectl get endpoints myservice**

```
labsuser@master:~$ kubectl get endpoints myservice
NAME           ENDPOINTS                                     AGE
myservice      192.168.232.208:80,192.168.47.146:80        2m39s
labsuser@master:~$
```

By following these steps, you have successfully demonstrated how to seamlessly integrate Kubernetes deployments with services to achieve scalable and accessible pod configurations.