

Lesson 05 Demo 04

Setting up an Ingress Controller with Transport Layer Security

Objective: To implement the transport layer security by deploying an Ingress rule to generate an SSL certificate

Tools required: kubeadm, kubectl, kubelet, and containerd

Prerequisites: A Kubernetes cluster (refer to Demo 01 from Lesson 01 for setting up a cluster)

Steps to be followed:

1. Deploy Ingress
2. Deploy HTTPD and OpenShift
3. Generate a self-signed SSL certificate and a TLS certificate
4. Verify the Ingress rule

Step 1: Deploy Ingress

1.1 Go to the terminal of your lab and execute the following command to deploy Ingress:

kubectl apply -f <https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.0/deploy/static/provider/cloud/deploy.yaml>

```
labsuser@master:~$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.1.0/deploy/static/provider/cloud/deploy.yaml
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
configmap/ingress-nginx-controller created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
service/ingress-nginx-controller-admission created
service/ingress-nginx-controller created
deployment.apps/ingress-nginx-controller created
ingressclass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
serviceaccount/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
job.batch/ingress-nginx-admission-create created
job.batch/ingress-nginx-admission-patch created
labsuser@master:~$
```

- 1.2 Verify all Ingress deployments, pods, and services using the following command:
kubectl get all -n ingress-nginx

```
labsuser@master:~$ kubectl get all -n ingress-nginx
NAME                                READY    STATUS      RESTARTS   AGE
pod/ingress-nginx-admission-create-s17lc    0/1     Completed   0           62s
pod/ingress-nginx-admission-patch-9h6fj     0/1     Completed   1           62s
pod/ingress-nginx-controller-6fcf745c45-dfb2m 1/1     Running     0           63s

NAME                                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)                                AGE
service/ingress-nginx-controller      LoadBalancer       10.98.39.108  <pending>      80:30380/TCP,443:32677/TCP            63s
service/ingress-nginx-controller-admission ClusterIP           10.101.55.45  <none>         443/TCP                                63s

NAME                                READY    UP-TO-DATE   AVAILABLE   AGE
deployment.apps/ingress-nginx-controller 1/1      1             1           63s

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.apps/ingress-nginx-controller-6fcf745c45 1          1          1        63s

NAME                                COMPLETIONS   DURATION   AGE
job.batch/ingress-nginx-admission-create    1/1           8s         62s
job.batch/ingress-nginx-admission-patch    1/1           9s         62s
labsuser@master:~$
```

- 1.3 List the created Ingress pods using the following command:
kubectl get pod -n ingress-nginx

```
labsuser@master:~$ kubectl get pod -n ingress-nginx
NAME                                READY    STATUS      RESTARTS   AGE
ingress-nginx-admission-create-s17lc    0/1     Completed   0           2m16s
ingress-nginx-admission-patch-9h6fj     0/1     Completed   1           2m16s
ingress-nginx-controller-6fcf745c45-dfb2m 1/1     Running     0           2m17s
labsuser@master:~$
```

Step 2: Deploy HTTPD and OpenShift

- 2.1 Deploy HTTPD and OpenShift deployments using the following commands:
kubectl create deployment myapp2 --image=docker.io/openshift/hello-openshift
kubectl create deployment myapp1 --image=docker.io/httpd
kubectl expose deployment myapp1 --port=80
kubectl expose deployment myapp2 --port=8080
kubectl get svc

```
labsuser@master:~$ kubectl create deployment myapp2 --image=docker.io/openshift/hello-openshift
deployment.apps/myapp2 created
labsuser@master:~$ kubectl create deployment myapp1 --image=docker.io/httpd
deployment.apps/myapp1 created
labsuser@master:~$ kubectl expose deployment myapp1 --port=80
service/myapp1 exposed
labsuser@master:~$ kubectl expose deployment myapp2 --port=8080
service/myapp2 exposed
labsuser@master:~$ kubectl get svc
NAME      TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)                                AGE
kubernetes ClusterIP   10.96.0.1      <none>         443/TCP                                13m
myapp1    ClusterIP   10.107.220.9   <none>         80/TCP                                 43s
myapp2    ClusterIP   10.105.105.22  <none>         8080/TCP                               26s
labsuser@master:~$
```

2.2 Create a directory for the Ingress TLS using the following commands:

```
mkdir ingress1
```

```
cd ingress1/
```

```
labsuser@master:~$ mkdir ingress1
labsuser@master:~$ cd ingress1
labsuser@master:~/ingress1$
```

Step 3: Generate a self-signed SSL certificate and a TLS certificate

3.1 Generate a self-signed SSL certificate using the following OpenSSL command:

```
openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout ingress.key-out
ingress.crt -subj "/CN=master.example.com/O=security"
```

[illegible]

3.2 Create a **tls-cert** certificate using the following command:

```
kubectl create secret tls tls-cert --key ingress.key --cert ingress.crt
```

```
labsuser@master:~/ingress$ kubectl create secret tls tls-cert --key ingress.key --cert ingress.crt
secret/tls-cert created
labsuser@master:~/ingress$
```

Step 4: Verify the Ingress rule

4.1 Create a YAML file for the Ingress rule using the following command:

vi rule.yaml

[illegible]

4.2 Add the following code within the **rule.yaml** file:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /$2
  name: rewrite
spec:
  tls:
  - hosts:
    - master.example.com
    secretName: tls-cert
  ingressClassName: nginx
  rules:
  - host: master.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: myapp1
            port:
              number: 80
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /$2
  name: rewrite
spec:
  tls:
  - hosts:
    - master.example.com
    secretName: tls-cert
  ingressClassName: nginx
  rules:
  - host: master.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: myapp1
            port:
              number: 80
```

4.3 Create and verify the Ingress rule using the following commands:

```
kubectl create -f rule.yaml
```

```
kubectl get ingress
```

```
labsuser@master:~/ingress1$ kubectl create -f rule.yaml
ingress.networking.k8s.io/rewrite created
labsuser@master:~/ingress1$ kubectl get ingress
NAME      CLASS      HOSTS          ADDRESS      PORTS      AGE
rewrite   nginx      master.example.com  80, 443      25s
```

4.4 Execute the following command to get the IP address associated with ens5:

```
ip a | grep ens5
```

```
labsuser@master:~/ingress1$ ip a | grep ens5
2: ens5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP group default qlen 1000
    inet 172.31.34.109/20 metric 100 brd 172.31.47.255 scope global dynamic ens5
labsuser@master:~/ingress1$
```

4.5 Execute the following command to list the IP hostnames and addresses for the localhost:

```
sudo vi /etc/hosts
```

```
labsuser@master:~/ingress1$ sudo vi /etc/hosts
labsuser@master:~/ingress1$
```

```
127.0.0.1 localhost

# The following lines are desirable for IPv6 capable hosts
::1 ip6-localhost ip6-loopback
fe00::0 ip6-localnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters
ff02::3 ip6-allhosts
```

4.6 List the running Ingress pods using the following commands:

```
kubectl get svc -n ingress-nginx
```

```
kubectl get pod -n ingress-nginx -o wide
```

```
kubectl get nodes -o wide
```

```
labsuser@master:~/ingress1$ kubectl get svc -n ingress-nginx
NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
ingress-nginx-controller            LoadBalancer  10.98.39.108   <pending>      80:30380/TCP,443:32677/TCP  33m
ingress-nginx-controller-admission  ClusterIP     10.101.55.45   <none>         443/TCP           33m
labsuser@master:~/ingress1$ kubectl get pod -n ingress-nginx -o wide
NAME                                READY    STATUS    RESTARTS   AGE   IP              NODE                                NOMINATED NODE   READINESS GATES
ingress-nginx-admission-create-s17lc 0/1      Completed 0           33m   192.168.47.130  worker-node-1.example.com          <none>           <none>
ingress-nginx-admission-patch-9h6fj   0/1      Completed 1           33m   192.168.47.129  worker-node-1.example.com          <none>           <none>
ingress-nginx-controller-6fcf745c45-df6zm 1/1      Running   0           33m   192.168.232.193 worker-node-2.example.com          <none>           <none>
labsuser@master:~/ingress1$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION   INTERNAL-IP    EXTERNAL-IP    OS-IMAGE             KERNEL-VERSION      CONTAINER-RUNTIME
master.example.com                   Ready    control-plane  40m   v1.28.2   172.31.34.109   <none>         Ubuntu 22.04.3 LTS   6.2.0-1012-aws     containerd://1.6.8
worker-node-1.example.com            Ready    <none>      35m   v1.28.2   172.31.24.250   <none>         Ubuntu 22.04.3 LTS   6.2.0-1012-aws     containerd://1.6.8
worker-node-2.example.com            Ready    <none>      35m   v1.28.2   172.31.18.47    <none>         Ubuntu 22.04.3 LTS   6.2.0-1012-aws     containerd://1.6.8
```

This shows the NodeIP and NodePort of the Ingress service.

4.7 Verify the generated certificate using the following commands:

kubectl get svc -n ingress-nginx

curl -kv https://master.example.com:31909/test

```
labsuser@master:~/ingress1$ kubectl get svc -n ingress-nginx
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP  PORT(S)                                AGE
ingress-nginx-controller            LoadBalancer        10.104.114.248   <pending>    80:31909/TCP,443:31999/TCP            5d19h
ingress-nginx-controller-admission ClusterIP            10.100.65.46     <none>       443/TCP                                5d19h
labsuser@master:~/ingress1$ curl -kv https://master.example.com:31909/test
* Trying 172.31.55.157:31909...
* TCP_NODELAY set
* Connected to master.example.com (172.31.55.157) port 31909 (#0)
* ALPN, offering h2
* ALPN, offering http/1.1
* successfully set certificate verify locations:
*   CAfile: /etc/ssl/certs/ca-certificates.crt
*   CApath: /etc/ssl/certs
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
```

By following these steps, you have successfully implemented the transport layer security by deploying an Ingress rule to generate an SSL certificate.